COACH ROOM WORKSHEET

Architecture Narrative Worksheet

A decision-led Senior-to-Staff narrative system for making architecture judgment, influence beyond authority, and bounded outcomes easy to inspect.

Lead with the expensive decision, connect every component to a constraint, and prove influence and consequence without inflating ownership.

VERSION 1.0

FORMAT COACH WORKSHEET + COMPLETED EXAMPLES

STATUS LEARNER READY

COACH-ROOM / ARCHITECTURE-NARRATIVE

MAKE THE LEVEL SIGNAL VISIBLE

An architecture story is a decision record spoken aloud

A strong Senior-to-Staff narrative lets the listener inspect the system constraint, the expensive decision, your authority boundary, the people you influenced, and the consequence that persisted after launch. The comparison below is a fictional teaching example; replace every fact, date, metric, and ownership claim with your own verified evidence.

CORE IDEA

Do not narrate every component equally. Lead with the choice that was costly to reverse, then use architecture, influence, and outcome evidence to prove why it mattered.

01	02	03	04
Pressure	**Decision**	**Influence**	**Consequence**
Name the user, operating, or business constraint that forced a decision.	State the options, the chosen path, the accepted cost, and what would reverse it.	Show whose behavior or plan had to change and how alignment was earned.	Connect the decision to a measured result, durable mechanism, and honest limit.

What the listener must be able to repeat

- ☐ The system problem and the non-negotiable constraint.
- ☐ The primary decision you personally shaped, not just the final architecture.
- ☐ The strongest rejected alternative and the trade-off you accepted.
- ☐ The stakeholder tension, decision owner, and influence move that changed the path.
- ☐ The baseline, observation window, result, and limit of the evidence.

Fictional micro-example: project tour versus staff-level narrative

Signal	Weak version	Decision-led version
Opening	We rebuilt the platform with several services.	A 40-minute recovery target could not survive the shared database failure domain.
Ownership	We decided to migrate.	I framed the decision, tested two alternatives, and brought the final owner an evidence-backed recommendation.
Trade-off	The new design scaled better.	We accepted duplicate processing during retry and paid the cost with idempotency plus reconciliation.
Outcome	Reliability improved a lot.	Recovery fell from 38 to 11 minutes across six observed incidents; low-volume tenants were not yet tested.

CHOOSE THE RIGHT STORY

Select for decision density, not project size

The best narrative contains a real fork in the road, observable ownership, cross-boundary coordination, and evidence after the change. A famous project with no decision tension is usually weaker.

Selection test	What qualifies	What weakens the story
Decision	Two credible paths existed and each carried a cost.	The answer was obvious or dictated before you entered.
Personal scope	You framed, recommended, tested, sequenced, or changed the decision.	Your role cannot be separated from the team summary.
Influence	Another team, leader, or function changed a commitment or behavior.	Only your direct reports needed instruction.
Consequence	A user, operating, delivery, or organizational result was observed.	Only implementation completion is available.
Truth boundary	Sources, dates, estimates, and unknowns can be stated precisely.	The story requires invented scale, causality, or ownership.

Candidate story comparison

FICTIONAL WORKED EXAMPLE - ADAPT THE METHOD, NOT THE FACTS

Candidate	Decision tension	Influence evidence	Outcome evidence	Choice
A: notification migration	Dual-write speed versus one source of truth	Three consuming teams changed adoption sequence	Six-week delivery and on-call trend	Primary
B: dashboard rewrite	Library choice after direction was set	One local team	Launch completed; adoption not measured	Backup

COVERAGE

Use a backup story

Keep a second story with a different decision type: reliability, migration, cost, platform adoption, or operating model. It protects against a probe that makes the primary story a poor fit.

BOUNDARY

Reject an unsafe story

Do not use confidential architecture, another person's decision, unverified impact, or a story whose evidence cannot be separated from later unrelated work.

BUILD THE DECISION SPINE

Make the irreversible choice easy to inspect

Treat the narrative as a compact architecture decision record. Supporting components belong only when they explain a constraint, a trade-off, a failure path, or the measured consequence.

01
Contract
Users, guarantee, scale, and decision deadline.

02
Constraint
The pressure that eliminated the default path.

03
Options
Two credible alternatives evaluated on the same criteria.

04
Choice
Decision, accepted cost, safeguards, and owner.

05
Reversal
Evidence or condition that would trigger a different path.

Layer	Question to answer	Evidence to carry	Red flag
Contract	What had to remain true for users and operators?	SLO, data rule, volume band, deadline	Starting with services and tools

Layer	Question to answer	Evidence to carry	Red flag
Constraint	Which limit made the choice non-trivial?	Failure history, cost, coupling, compliance, migration window	Calling every requirement critical
Options	Why were both paths plausible?	Same criteria, test result, prototype, operating cost	A straw-man rejected option
Choice	What did you choose and pay for?	Accepted risk, mitigation, decision owner, sequence	Only describing benefits
Reversal	When would the decision stop being correct?	Threshold, new fact, review date, exit path	Pretending the choice is permanent

DECISION QUALITY

Rejected alternative

Represent the strongest alternative fairly. Name what it optimized, why it lost under the current constraints, and the future condition under which it could win.

TRADE-OFF

Accepted cost

Every credible decision spends something: latency, complexity, migration time, consistency, money, or team attention. State the cost before the interviewer has to ask.

CARRY ARCHITECTURE EVIDENCE

Connect each component to a reason

A diagram is useful only when the learner can explain the system before and after the decision, the invariant that protects correctness, and what happens when the happy path breaks.

Architecture evidence	Minimum useful explanation	Proof or artifact
Before state	Where state lived, the critical flow, and the recurring failure or delivery cost.	Incident timeline, delivery lead time, dependency map
After state	What changed in ownership, data flow, boundary, or operating responsibility.	Decision record, migration plan, interface contract
Invariant	The rule that must remain true across retries, races, partial failure, or migration.	Test, reconciliation rule, version check, idempotency key
Failure path	Detection, containment, recovery, and the unsafe failure you explicitly rejected.	Alert, rollback, replay, dead-letter or repair path
Rollout	How risk was bounded while evidence accumulated.	Shadow read, cohort, dual run, feature flag, stop threshold
Operating loop	Who owns the signal and what decision the metric triggers.	Dashboard owner, review cadence, runbook action

NARRATION RULE

Use one sentence per component: the constraint it answers, the failure it contains, or the evidence it enables. If a box serves none of those purposes, remove it from the spoken story.

A before-and-after explanation

FICTIONAL WORKED EXAMPLE - ADAPT THE METHOD, NOT THE FACTS

Moment	Spoken evidence	Why it earns signal
Before	A single scheduler owned priority and delivery state, so one noisy tenant delayed every queue.	Names the coupling and user consequence.
Decision	Partition admission by tenant class, keep delivery state durable, and preserve at-least-once processing with idempotent consumers.	States boundary, guarantee, and accepted retry behavior.
Failure	If a worker dies after send, the durable attempt is retried; duplicate risk is bounded by the idempotency contract and reconciliation.	Explains uncertainty without claiming impossible ownership.
Rollout	Shadow the new scheduler, compare age and duplicate rates, then move three cohorts with a rollback threshold.	Shows reversible evidence gathering.

SHOW INFLUENCE WITHOUT INFLATING AUTHORITY

Trace how a decision moved across boundaries

Staff-level influence is not a list of meetings. It is evidence that you understood competing incentives, changed the decision process, and left a mechanism that continued working without personal heroics.

01
Map
Identify decision owner, affected teams, incentives, and cost of the ask.

02
Listen
State the strongest objection in terms the other stakeholder would accept.

03
Reduce
Offer the smallest reversible commitment that can produce decision evidence.

04
Durabilize
Leave an interface, review, ownership rule, or metric after alignment.

Influence map

FICTIONAL WORKED EXAMPLE - ADAPT THE METHOD, NOT THE FACTS

Actor	Competing constraint	Your move	Decision evidence
Payments team	Could not absorb a breaking event contract during quarter close.	Proposed a versioned adapter and a two-cohort migration instead of a single cutover.	Committed to cohort one after replay tests passed.
Platform owner	Wanted one reusable scheduler rather than tenant-specific forks.	Made tenant class a policy input and documented the extension boundary.	Approved the shared contract and ownership split.
Support lead	Needed predictable escalation when delivery age crossed the customer promise.	Added an age-based signal and named the responding owner in the runbook.	Used the signal during the first production incident.
Director	Needed a decision before the staffing window closed.	Presented options, accepted cost, open risk, and the reversal threshold in one page.	Selected staged migration with a dated review.

Common influence mistakes

+ Saying I convinced everyone without naming the legitimate objection or decision owner.
+ Using meeting count as impact instead of showing which commitment changed.
+ Claiming the team's execution as personal ownership.
+ Treating escalation as failure instead of explaining why the decision needed the right owner.
+ Ending at agreement without the durable mechanism that kept teams aligned.

PROVE THE CONSEQUENCE

Use bounded evidence instead of impact theater

A result is credible when the listener can see the baseline, the change, the observation window, the measured movement, and what the data does not prove.

Evidence layer	Question	Good answer pattern
Baseline	What was true before the decision?	Metric or recurring observable behavior, with source and date.
Change	What action could plausibly affect it?	Specific architecture, rollout, ownership, or operating change.
Window	When and for how long was it observed?	A stated cohort and period, not immediately after launch.
Result	What moved and by how much?	Exact measure when stored; qualitative evidence labeled when not.
Limit	What remains uncertain or confounded?	Missing cohort, concurrent change, small sample, or proxy limitation.
Durability	What continued after you stepped away?	Ownership, review, interface, runbook, or automated guardrail.

FICTIONAL WORKED EXAMPLE - ADAPT THE METHOD, NOT THE FACTS

Baseline	Change	Window	Observed result	Honest limit
Priority queue age exceeded 18 minutes in four of seven peak periods.	Tenant-aware admission, durable attempt state, and an age alert with named owner.	Six weeks; three migrated cohorts.	Peak age stayed below six minutes; two incidents recovered without manual queue surgery.	Low-volume cohorts and seasonal peak behavior were not yet observed.

USE CAREFULLY

Leading signal

A rehearsal, rollout, adoption, latency, or error signal can show movement early. Label it as leading evidence; do not present it as the final business outcome.

STRONGER PROOF

Lagging outcome

Customer, reliability, cost, or delivery impact needs enough observation time and a plausible connection to the change. Name concurrent factors when they exist.

ASSEMBLE THE SPOKEN NARRATIVE

Use a decision-first spine at three depths

Prepare one truthful source narrative, then rehearse a 90-second summary and a four-minute version. Do not memorize separate stories that drift apart under follow-up.

Clock	Purpose	Content
0:00-0:30	Orient	User or operating problem, scope, personal role, and the decision deadline.
0:30-1:15	Create tension	Constraint, two credible options, and why the default path was unsafe or insufficient.
1:15-2:45	Show judgment	Chosen path, accepted cost, safeguards, rollout, failure path, and reversal condition.
2:45-3:35	Show influence	Competing constraint, your move, changed commitment, final decision owner, durable mechanism.
3:35-4:15	Prove consequence	Baseline, observation window, result, limit, and what you learned or would change.
4:15-close	Invite depth	One sentence that names the richest technical or influence follow-up available.

Fictional four-minute spine

FICTIONAL WORKED EXAMPLE - ADAPT THE METHOD, NOT THE FACTS

Beat	Illustrative line
Pressure	Our shared notification scheduler breached the 15-minute delivery promise when one large tenant burst, and three teams depended on the same event contract.
Decision	I compared a larger shared queue with tenant-aware admission. I recommended the second path because it bounded noisy-neighbor risk, while accepting policy and reconciliation complexity.
Architecture	We kept durable attempt state, at-least-once delivery, an idempotency contract, and a staged cohort migration with age and duplicate thresholds.
Influence	Payments could not take a breaking change. I reframed the migration around a versioned adapter and replay proof; the platform owner retained the final interface decision.
Outcome	Across six weeks and three cohorts, peak age stayed under six minutes and two incidents recovered through the new runbook. Low-volume cohorts remained an open evidence gap.

FIRST ANSWER

90-second compression

Keep pressure, primary decision, strongest trade-off, one influence move, one bounded outcome, and the limit. Remove component detail before removing the decision.

FOLLOW-UP

Deep-dive expansion

Expand only the layer the interviewer probes: option criteria, invariant, failure timeline, stakeholder objection, rollout evidence, or causal limit.

REHEARSE FOR INSPECTION

Use probes to test ownership, judgment, and truth boundaries

A credible narrative survives follow-up without changing facts. Record the sentence where evidence becomes inference, the question that exposes a weak layer, and one correction for the next run.

Probe	What it tests	Evidence-ready response
Why was this the decision to lead with?	Prioritization	Cost of delay, reversibility, user or operating risk, and decision deadline.
Why not the strongest alternative?	Trade-off quality	Shared criteria, what the alternative optimized, why it lost now, and reversal condition.
What was specifically yours?	Ownership precision	Your framing, analysis, recommendation, sequencing, or mechanism; final owner and team work separated.
Who disagreed and why?	Influence maturity	A fair competing constraint, the smallest commitment, evidence produced, and changed decision.
What failed during rollout?	Operating depth	Detection, containment, recovery, user impact, and the safeguard changed afterward.
What did the metric not prove?	Evidence honesty	Observation window, sample, proxy, concurrent change, and remaining unknown.

0-2 evidence rubric

Dimension	0 - absent or unsafe	1 - present but weak	2 - inspectable
Decision	No real fork	Choice named; criteria vague	Options, criteria, cost, owner, reversal
Architecture	Component tour	Flow clear; failure thin	Invariant, failure, rollout, operating signal
Ownership	Team-only claims	Personal role broad	Personal decision separated from team and owner
Influence	Agreement asserted	Stakeholders named	Constraint, move, changed commitment, durable mechanism
Outcome	Unbounded claim	Metric without context	Baseline, window, result, limit, durability
Delivery	Chronology without thesis	Decision arrives late	Decision-first, concise, probe-ready

COACH HANDOFF

Bring two story cards, one recorded four-minute run, the lowest rubric row with a timestamped quote, and one corrected sentence. That is enough evidence for a focused calibration session.

Completion check

- [ ] The primary decision appears in the opening minute.
- [ ] The rejected alternative is credible and the accepted cost is explicit.
- [ ] Personal scope, team execution, and final authority are separated.
- [ ] Influence ends in a changed commitment and a durable mechanism.
- [ ] The outcome includes a baseline, window, result, and honest limit.
- [ ] The 90-second and four-minute versions preserve the same facts.

Use after every rehearsal

01
Observe
Save one exact quote or timestamp from the lowest-scoring dimension.

02
Correct
Rewrite one sentence or reorder one decision; do not rebuild the whole story.

03
Transfer
Retest the correction with a fresh probe and stop when it appears without rescue.

NEXT ACTION

Choose two truthful architecture stories now. Score both, select the decision-dense primary story, and record one four-minute run for coach review.