OPERATIONS WORKBOOK

Evidence Comparison and Correction Workbook

A three-attempt review system for isolating one persistent gap, choosing a correction, and running a comparable re-test.

Find the repeated decision moment behind the score, change one behavior, and require evidence of transfer.

VERSION 1.0

FORMAT GUIDE + WORKSHEET

STATUS LEARNER READY

OPERATIONS / EVIDENCE-COMPARISON-CORRECTION

COMPARE WITH DISCIPLINE

Find the repeated decision, not the lowest number

Three attempts become useful only when the prompt shape, conditions, and rubric are comparable enough to reveal a recurring behavior.

CORE IDEA

A persistent gap is a repeated observable moment under comparable conditions. One awkward answer is a hypothesis; three low totals without evidence are still only three numbers.

01	02	03	04
Normalize	**Quote**	**Cluster**	**Correct**
Use the same dimensions and record prompt difficulty and conditions.	Capture the exact moment the decision became unclear or late.	Look for repeated cues, causes, and downstream consequences.	Choose one front-loaded behavior and define the re-test.

PERSISTENT	ONE-OFF
Signal	**Noise**
The same behavior appears in similar decision moments and affects correctness, clarity, trade-offs, or recovery.	The prompt, time, assistance, fatigue, rubric, or domain changed enough that the result is not directly comparable.

Before comparing

- [ ] Prompt type and difficulty are recorded.
- [ ] Time, tools, and assistance are visible.
- [ ] The same rubric dimensions were used.
- [ ] Evidence includes quotes, timestamps, code/tests, diagrams, or reviewer notes.

USE ONE SHARED RUBRIC

Compare behavior across interview types

The wording changes by domain, but the comparison should still expose framing, decision quality, boundary handling, and communication.

Dimension	Coding evidence	System-design evidence	Behavioral evidence
Framing	Contract, examples, constraints	Requirements, quality attributes, non-goals	Question, context, stakes
Decision	Pattern and state choice	Architecture/data choice	Your choice, not only team action
Reason	Invariant and complexity proof	Requirement and trade-off	Constraint, alternative, judgment
Boundary	Empty, duplicate, off-by-one	Failure, overload, stale state	Conflict, uncertainty, follow-up
Evidence	Dry run, tests, correct output	Estimate, metric, recovery behavior	Measured result and durable change
Communication	Decisions audible at the right time	Flow remains traceable as design expands	Scope is clear without excessive setup

SCORING DISCIPLINE

Use the same 0 / 1 / 2 evidence scale from the Technical Baseline Scorecard, but compare the notes behind each score. The number is an index into evidence, not the conclusion.

A useful quote

Record enough surrounding context to understand the decision. 'I forgot consistency' is weaker evidence than: 'I selected asynchronous replication, then described read-your-writes only after the interviewer asked what a user sees after an update.'

FIND THE REPEATED MOMENT

Trace symptom, cause, and downstream consequence

Do not stop at the visible mistake. Find the earliest decision whose absence made later parts of the answer harder to recover.

Repeated symptom	Possible earlier cause	Downstream consequence	Correction candidate
Edge cases appear only after prompting	The input contract and invariant were never stated	Implementation changes late; testing becomes random	Name contract and one boundary before coding
Design becomes a list of components	Requirements and first bottleneck were not used as decision anchors	Trade-offs sound generic; failures are bolted on	State requirement -> decision -> cost for each major box
Ownership sounds team-wide	The speaker never named the personal decision boundary	Level and judgment are hard to calibrate	Open with context, then 'My decision was...'
Follow-up answers contradict the opening	Assumptions were implicit and changed silently	The interviewer cannot tell whether the design adapted coherently	Label the assumption before changing the decision

Root-cause questions

+ What cue was present before the gap appeared?
+ Which decision should have been made at that moment?
+ Did the same absence create problems in more than one attempt?
+ Would correcting it improve several downstream dimensions?
+ Can the behavior be observed in the next comparable attempt?

CHOOSE THE SMALLEST CORRECTION

Prefer the earliest observable behavior with broad downstream impact. Do not replace a correction with a new syllabus.

WORKED EXAMPLE

From three attempts to one correction

The fictional evidence below shows how an apparent knowledge gap can actually be a decision-order problem.

Attempt	Prompt	Observed moment	Consequence
A	Notification system	Retry policy added after the queue diagram	Duplicate effects and overload path stayed vague
B	File processing	Consumer failure discussed only after a follow-up	No dead-letter or poison-message plan
C	Metrics pipeline	Delivery semantics named after storage choice	Reconciliation and idempotency contradicted earlier claims

TOO BROAD

Weak conclusion

I need to learn queues again. This is broad, hides what was already known, and gives no moment to test.

SPECIFIC

Evidence-led conclusion

I understand queue components, but delivery and retry semantics appear after architecture. Front-load them immediately after requirements.

CORRECTION

Before drawing components, state: critical work is at-least-once; consumers are idempotent; retries are bounded with dead letters; oldest-event age is the recovery signal.

Re-test condition

Use a nearby asynchronous-workflow prompt, the same 35-minute structure, and the same rubric. Success means delivery/retry behavior appears before architecture and remains coherent during one failure follow-up.

WORKSHEET 1

Compare the last three attempts

Normalize the attempt conditions first. Then score the same six dimensions and attach a short evidence note to every score.

Context	Attempt A	Attempt B	Attempt C
Date / prompt			
Type / difficulty			
Time / tools / assistance			
Rubric used			

Shared dimension	Attempt A: 0/1/2 + evidence	Attempt B: 0/1/2 + evidence	Attempt C: 0/1/2 + evidence
Framing			
Decision			
Reason			
Boundary			
Evidence			
Communication			

WORKSHEET 1B

Turn the score grid into evidence

Totals can hide the moment that matters. Use artifacts and repeated behavior to explain what is stable, what is weak, and what is still uncertain.

Synthesis	Evidence across Attempts A / B / C
Strongest repeated behavior	
First repeated weak decision moment	
Cue present immediately before it	
Quote / timestamp / artifact references	
What is noise or still uncertain?	

COMPARISON RULE

A score without an evidence note is incomplete. If prompt type, difficulty, assistance, or time changed enough to explain the difference, mark the comparison as uncertain instead of forcing a trend.

WORKSHEET 2

Select one correction

Write one line for each step. If the chain cannot be stated clearly, collect better evidence before choosing the correction.

Step	Your evidence
Repeated observable moment	
Cue present before it	
Earliest missing decision	
Downstream consequence	
One correction sentence	
Evidence that would show improvement	

CORRECTION SENTENCE TEMPLATE

When I notice [cue], I will [observable behavior] before [decision point], so [quality or failure consequence] stays explicit.

Keep / change

COMPARABLE

Keep stable

Prompt shape, time box, tools, assistance, rubric dimensions, and scoring evidence.

ATTRIBUTION

Change once

The one intended behavior. Do not add several new techniques during the same re-test.

WORKSHEET 3

Run the targeted re-test

Use one 35-minute block: spend 28 minutes replaying the original decision window, then use a seven-minute adjacent transfer probe. The probe is not a second full attempt.

Re-test setup	Your plan
Phase 1: original lowest-score scenario	
Same prompt conditions / decision-window time box / tools	
Correction cue	
Where the behavior must appear	
Primary success evidence	
Phase 2: 7-minute adjacent transfer probe	

Result	Evidence / quote / timestamp	Next decision
Phase 1: same scenario		Keep / adjust / gather evidence
Phase 2: brief transfer probe		Stable / prompt-specific / unclear

STOP CONDITION

Return to the main roadmap when the correction appears unaided in the comparable re-test and at least once in a nearby prompt. Do not rehearse until the answer becomes a script.

FINISH THE REVIEW

Make the next test smaller and clearer

The review is complete when one repeated gap has evidence, one correction has a visible moment, and the re-test conditions are recorded.

01	02	03	04
Compare	**Quote**	**Correct**	**Re-test**
Normalize conditions and use one rubric.	Capture repeated moments, not vague impressions.	Choose the earliest behavior with downstream impact.	Change one variable and check transfer.

Common mistakes

+ Picking the lowest score without reading the evidence behind it.
+ Comparing prompts or conditions that are not meaningfully comparable.
+ Calling a single bad moment persistent because it felt uncomfortable.
+ Choosing several corrections and losing attribution.
+ Repeating the identical answer until it becomes memorized rather than transferable.

YOUR NEXT ACTION

Complete the comparison grid, write one correction sentence, and schedule one comparable re-test plus one adjacent transfer check.

Completion check

- [ ] Three attempts use the same dimensions and include evidence.
- [ ] A repeated decision moment is separated from noise.
- [ ] One correction is attached to a cue and decision point.
- [ ] The comparable re-test and transfer check are defined.